

Average query latency vs. corpus size

